



Managing a Database Migration Project

Blair Layton

Business Development Manager, Database, Amazon Web Services



ย้ายเลย! ย้ายข้อมูลไปยัง AWS (ระดับ 200): จัดการโครงการย้ายข้อมูล DB ด้วยวิธีที่ดีที่สุด

Surawut Phornthabthong

Solutions Architect, Amazon Web Services

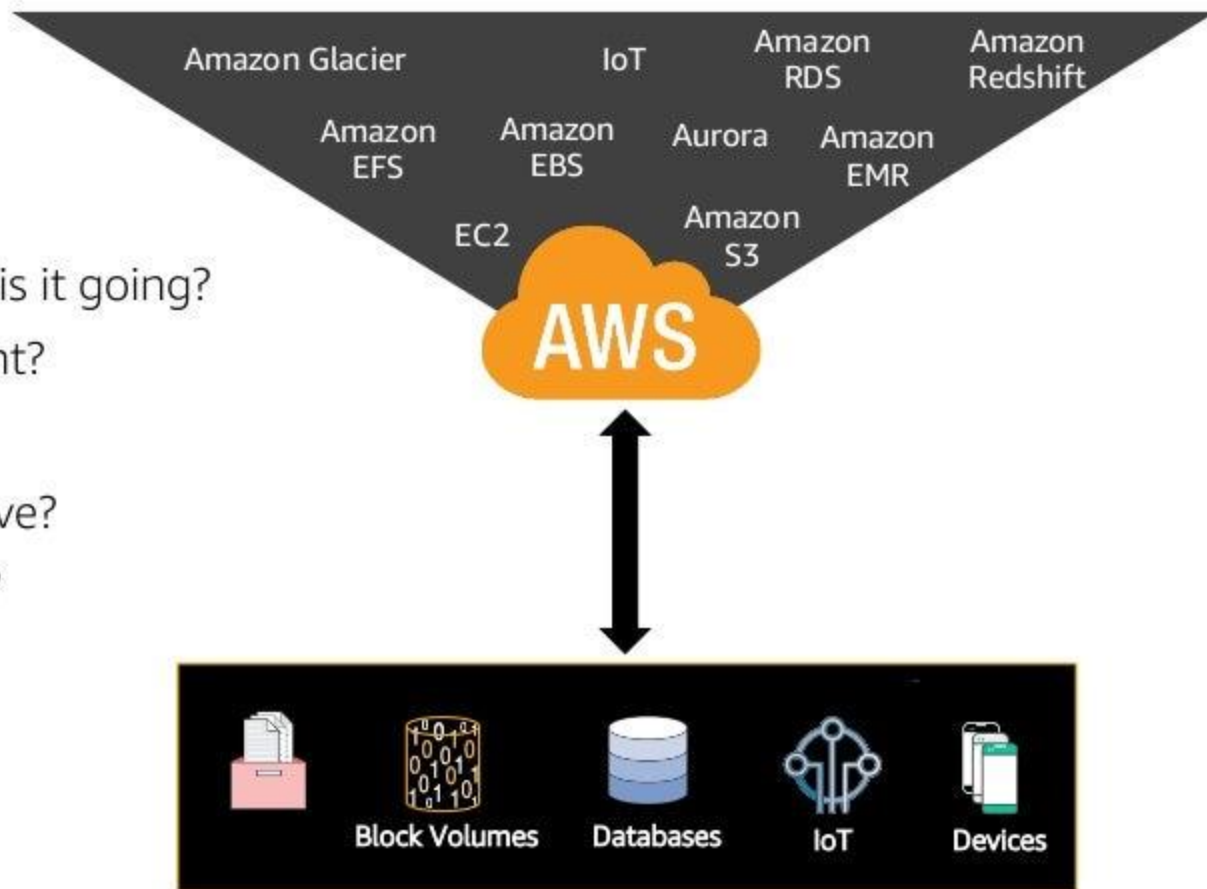
What to Expect from the Session

- Database Migration Context
- Database Migration Tools
- Introduction to the AWS Migration Framework
- Database Migration Effort
- Customer References
- Next Steps

Database Migration Context

Migrating Data: Five Key Questions

- 1) What kind of data is it, and where is it going?
- 2) One-time or continuous movement?
- 3) One-way or bidirectional access?
- 4) How much data & time do you have?
- 5) How might your WAN be a factor?



Amazon RDS Engines

Aurora



Open source



Commercial

ORACLE™

Microsoft SQL Server

Amazon DynamoDB

Fast and flexible NoSQL database service for any scale

Highly scalable



Automatic scaling to hundreds of terabytes of data that serve millions of requests per second

Fast, consistent performance



Consistent single-digit millisecond latency; DAX in-memory performance reduces response times to microseconds

Fully managed



Automatic provisioning, infrastructure management, scaling, and configuration with zero downtime

Business critical reliability



Data is replicated across fault-tolerant Availability Zones, with fine-grained access control

Migration Was Costly, Complex, & Slow

✗ Required commercial migration & replication software



✗ Caused long application downtime



✗ Was complex to set up & manage



✗ Required DB-specific application code

1001001
0101001
1000100

Database Migration Tools

AWS Database Migration Service (AWS DMS)

DMS migrates databases to AWS easily and securely with minimal downtime. It can migrate your data to and from most widely used commercial and open-source databases.



PostgreSQL

ORACLE



Amazon Aurora



MariaDB



mongoDB.



DynamoDB



AMAZON REDSHIFT



MySQL



Microsoft
SQL Server



Microsoft
SQL Azure

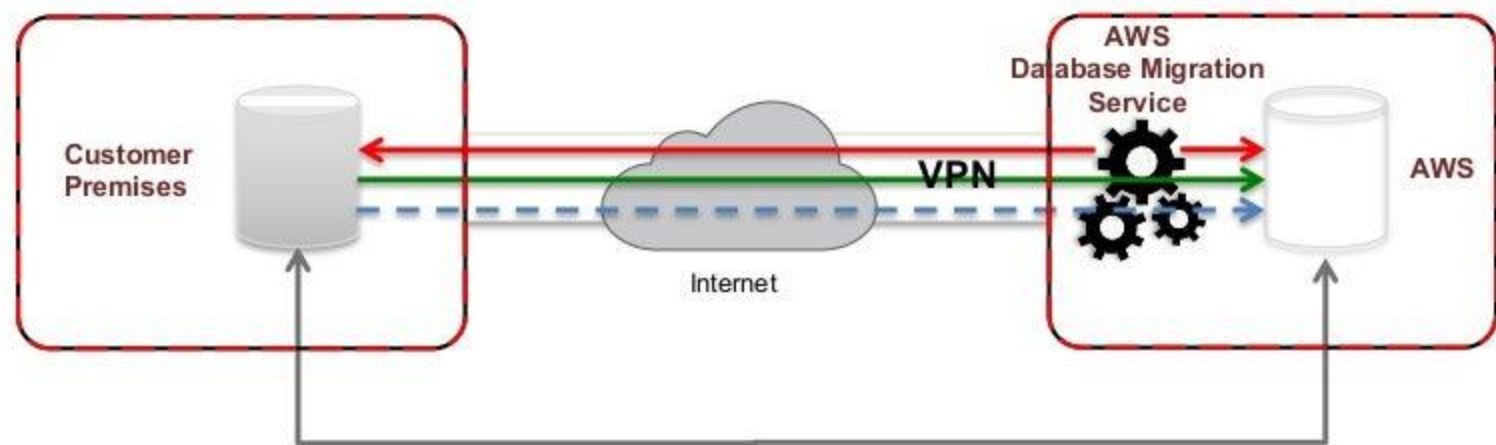


SAP ASE



S3 Bucket

Keep Your Apps Running During the Migration



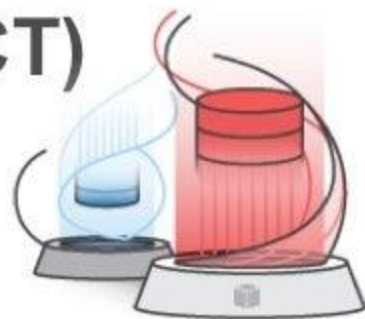
Start a replication instance
Connect to source and target
databases
Select tables, schemas, or
databases



Let AWS DMS create tables,
load data, and keep them in
sync
Switch applications over to the
target at your convenience

AWS Schema Conversion Tool (AWS SCT)

SCT helps automate many database schema and code conversion tasks when migrating between database engines or data warehouse engines



When to use DMS and SCT?

Tools for Migration Project Phases

Phase	Service/Tool	Notes
Assessment	AWS Schema Conversion Tool	Reports on the database objects, complexity and types of migration issues
Schema Migration	AWS Schema Conversion Tool	Copies a schema or migrates a schema depending on whether it is a homogeneous or heterogeneous migrations
Data Migration	AWS Database Migration Service AWS Schema Conversion Tool	Bulk load and change data capture (CDC) options Extraction and load for large data warehouses, including AWS Snowball integration
Application Migration	AWS Schema Conversion Tool	SQL statement migration in application code
Data Validation	AWS Database Migration Service	Ensure data is the same on source and target
Functional Testing	Various Tools on Marketplace	Ensure the application runs as intended
Performance Testing	Various Tools on Marketplace	Ensure the application performance as intended

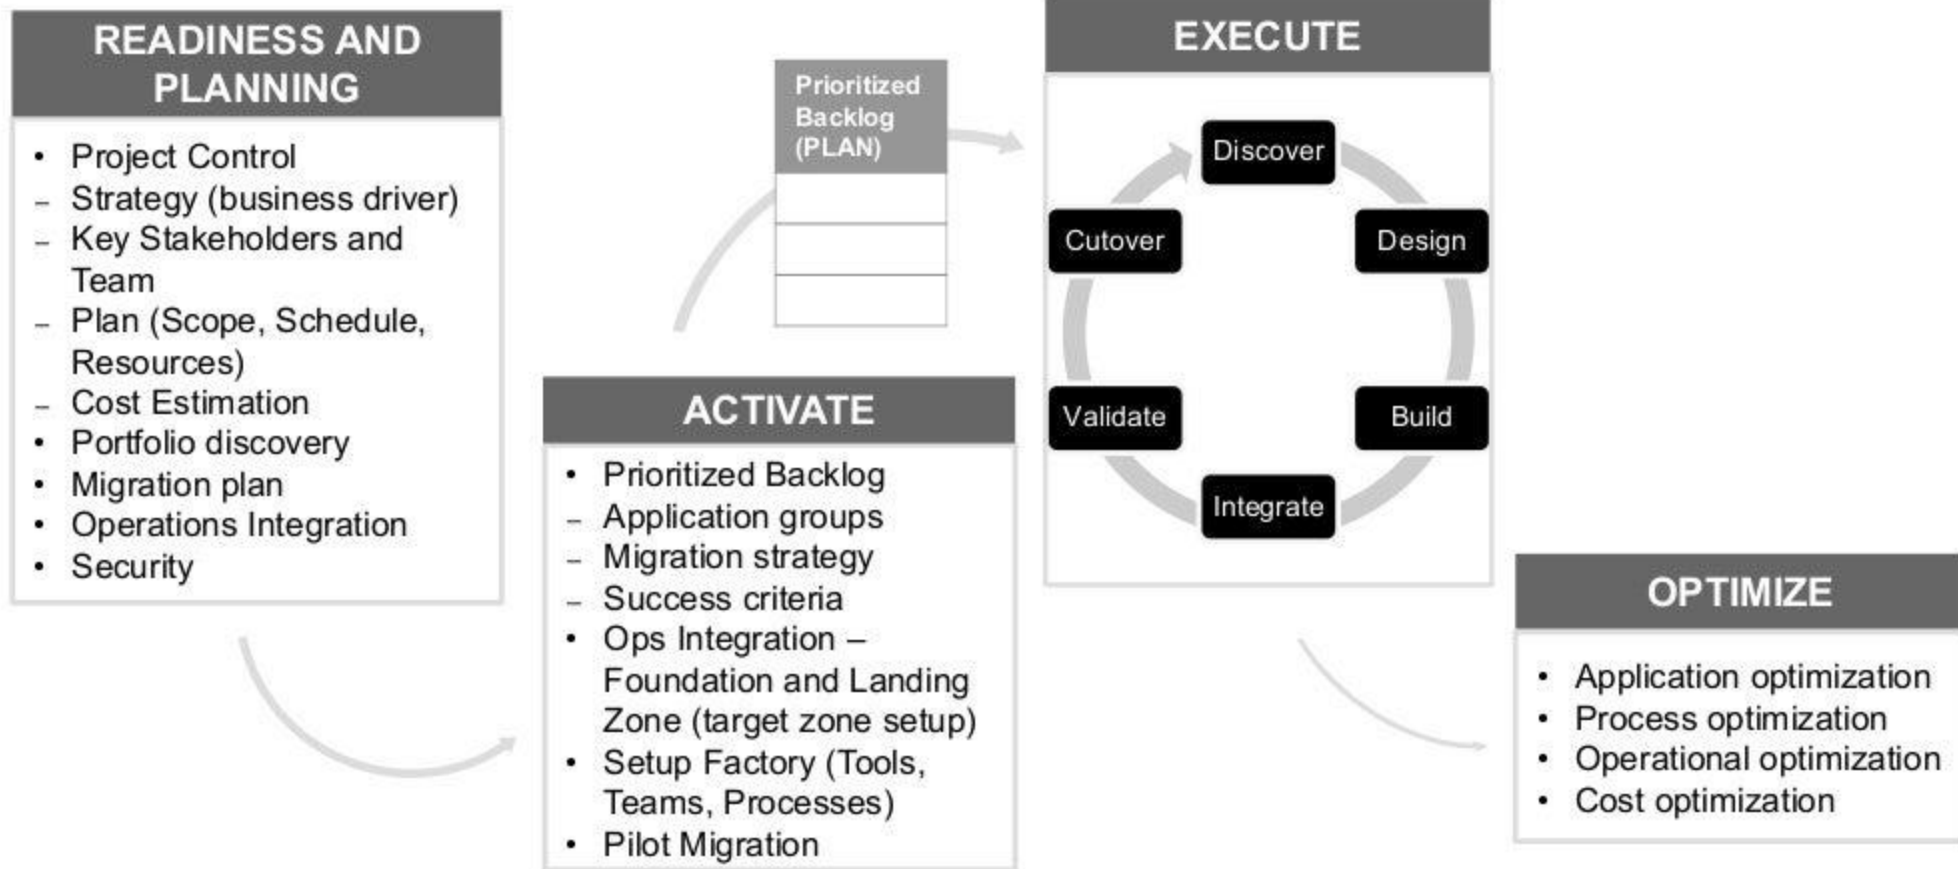
Tools for Migration Scenarios

Scenario	Example	Recommendation
Homogeneous migration to the same database version and edition	Migration of Oracle Database 11gR2 Enterprise Edition from on-premise to EC2	Use the native replication technology to create a standby database and then failover to the standby database
Homogeneous migration to a different version	Migration of MySQL 5.5 to MySQL 5.7	AWS Schema Conversion Tool and AWS Database Migration Service
Homogeneous migration to a different edition	Migration of SQL Server Enterprise Edition to Standard Edition	AWS Schema Conversion Tool and AWS Database Migration Service
Heterogeneous migration	Migration from Oracle Database to PostgreSQL	AWS Schema Conversion Tool and AWS Database Migration Service

**That's the tools, but how to
manage migration projects?**

Introducing the AWS Migration Framework

AWS Migration Framework



AWS Migration Framework- Readiness and Planning

AWS Migration Framework - Readiness & Planning

READINESS AND PLANNING

- **Project Control**
 - **Strategy (business driver)**
 - **Key Stakeholders and Team**
 - **Plan (Scope, Schedule, Resources)**
 - **Cost Estimation**
- Portfolio discovery
- Migration plan
- Operations Integration
- Security

Project Control focuses on ensuring there is a migration strategy in place that is **supported by key stakeholders** in the organisation. Additionally, we look at **defining the team** that will carry out the work, with associated **timelines** and **cost estimations**.

Sample decision points:

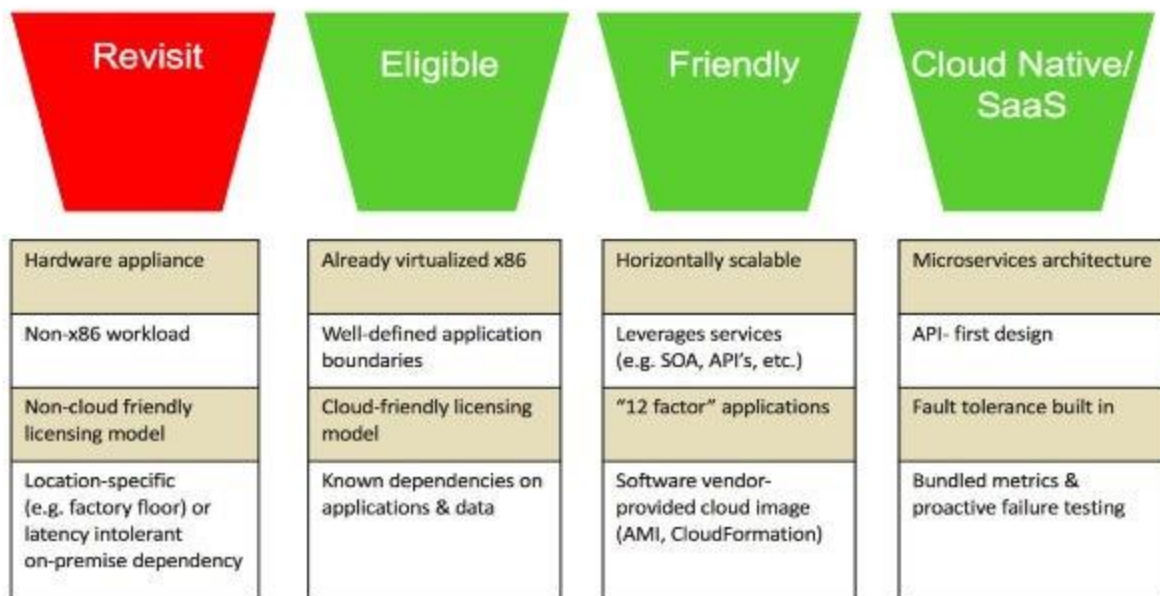
- Who is the executive sponsor?
- Are there any compelling events that will affect the migration strategy?
- Do we have the right resources? How are they organized?
- What are the timeframes we are working with?
- Do we have the necessary budget?

AWS Migration Framework - Readiness & Planning

READINESS AND PLANNING

- Project Control
 - Strategy (business driver)
 - Key Stakeholders and Team
 - Plan (Scope, Schedule, Resources)
 - Cost Estimation
- **Portfolio discovery**
- Migration plan
- Operations Integration
- Security

Quickly understand which applications are Cloud **Eligible**, Cloud **Friendly**, or Cloud **Native** and then execute a dive deep analysis on just that subset of applications.



*Friendly includes Eligible characteristics.

**Native/SaaS includes both Friendly and Eligible characteristics.

Application Assessment

- Business driver and intended ROI?
- Migration sponsor (business owner, C-level)?
- ISV application? Does the ISV support the target?
- Maintenance window for the migration?
- Design documentation?
- Original developers/DBAs still available?

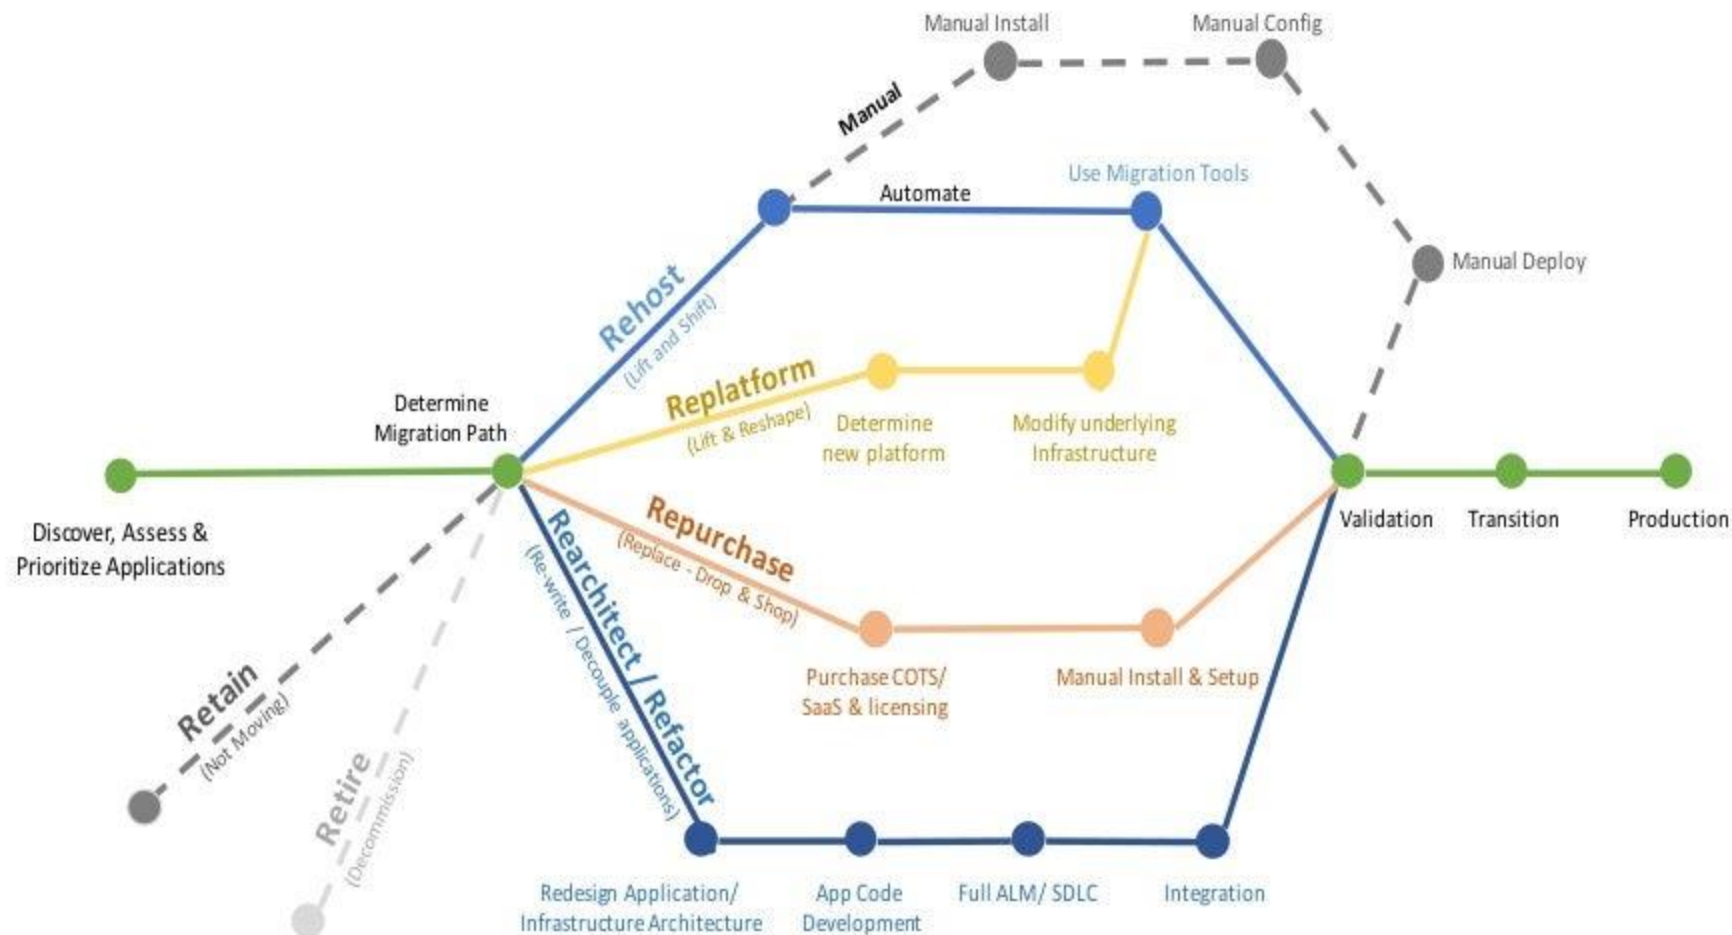
Database Assessment

- How many database objects (tables, triggers, SPs, users, etc.)?
- How much data?
- Complexity of the SPs and triggers?
- Proprietary DB features?
- Non-standard or custom data types?
- Character set conversions?
- Time zone or UTC?
- User authentication method?
- Licensing mechanism (cores, users, ULA etc.)

Application Technical Assessment

- Database Access:
 - SQL statements throughout the code?
 - Calls to a data abstraction layer?
 - API calls?
- ANSI SQL used where possible?
- SQL complexity, e.g. analytics with many joins or simple CRUD?
- Number of lines of SQL code?
- Application access, e.g. LDAP, DB Users, etc.

The 6Rs of Migration Planning



AWS Migration Framework - Activate

AWS Migration Framework - Activate

ACTIVATE

- Prioritized Backlog
 - Application groups
 - Migration strategy
 - Success criteria
 - Ops Integration – Foundation and Landing Zone (target zone setup)
 - Setup Factory (Tools, Teams, Processes)
 - Pilot Migration
- Determine your application priorities and group integrated applications together
 - Outline the success criteria for each application migration
 - Create your AWS landing zone (accounts, VPC, subnets, IAM roles, VPN/Direct Connect, etc.)
 - Configure DMS, SCT and other migration tools
 - Team creation
 - POC/pilot

Building a Migration Team

Application architect/developer: Application expert who can identify what components are important, complex, redundant, etc.

Source DBA: Knows the database design, schema, features used and what must be migrated to the target.

Target DBA: An expert in the target database to help map features from the source DB with the Source DBA.

AWS Solution Architect: Determines the correct target architecture in AWS and is familiar with DMS/SCT.

Application/Database Developers: Customer and/or partner resources to migrate the stored procedures, triggers and application code.

Hiring and Developing Talent

New skills are needed for the target DB and often AWS if migrating from on-premises

Develop **training plans** for existing employees

Hire in required skills if necessary

Retrain, redeploy or make people **redundant** who's skills are no longer relevant

Pilot/POC

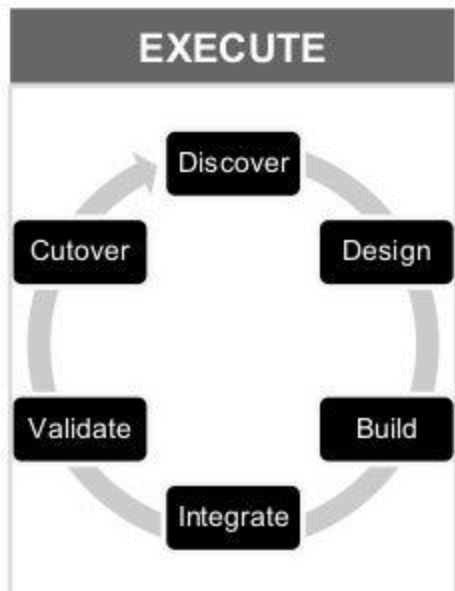
Choose a **reasonably complex** module/component to migrate to **validate your assumptions** in the **Activate** phase

You should:

- Obtain more accurate migration assessments
- Determine what can be **automated**
- Learn how the migration tools behave (limitations, bugs, improvements needed)
- Learn what **skills are missing** from your team

AWS Migration Framework - Execute

AWS Migration Framework - Execute



- Always have a back up plan!
- Execute according to lessons from Pilot/POC
- Typically the same amount of time to migrate the DB as to migrate the application (assuming DAL)
- Determine how you will cutover
 - **Parallel run**: expensive and difficult
 - **Minimal downtime**: DMS+CDC
 - **Large maintenance window**: application and data verification before go live

AWS Migration Framework - Optimize

AWS Migration Framework - Optimize

OPTIMIZE

- Application optimization
 - Process optimization
 - Operational optimization
 - Cost optimization
- DMS instance and task optimization
 - Database tuning
 - Database instance right sizing
 - Application tuning
 - Application instance right sizing
 - Use EC2 and RDS stop/start to optimize costs
 - Purchase reserved instances and use spot instances
 - Look for contention and evaluate caching, NoSQL, federation and adoption of a microservices strategy
 - Perform HA/DR scenarios and optimize the use of AWS managed services to help, e.g. RDS MAZ, Auto-scaling

Migration Effort

Database migration – multi phase process

Phase	Description	Automation	Effort (%)
1	Assessment	SCT	2
2	Database Schema Conversion	SCT/DMS	14
3	Application Conversion/Remediation	SCT	25
4	Scripts Conversion	SCT	7
5	Integration with 3 rd party applications		3
6	Data Migration	DMS	4
7	Functional testing of the entire system		29
8	Performance tuning	SCT	2
9	Integration and deployment		7
10	Training and knowledge		2
11	Documentation and version control		2
12	Post production support		3

Database Migration Process

STEP 1: Conversion

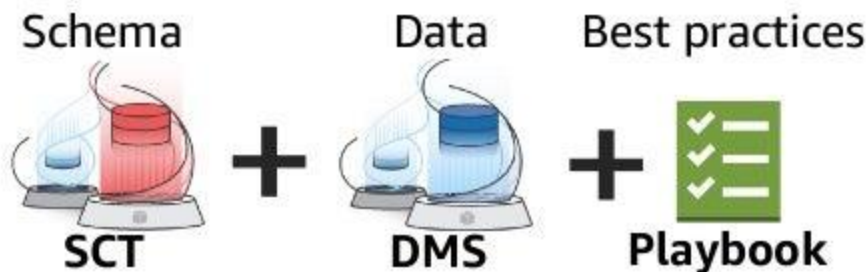


STEP 2: Data migration



Oracle to Aurora Migration Playbook

- Topic-by-topic overview of Oracle to Aurora PostgreSQL migrations and “hands-on” best practices
- How to migrate from proprietary features and the different database objects
- Migration best practices



	Oracle Feature	PostgreSQL Feature	Compatibility
Link	Index Organized Tables (IOTs)	PostgreSQL “Cluster” Tables	Yes*
Link	Common Data Types	Common Data Types	Yes
Link	Table Constraints	Table Constraints	Yes
Link	Table Partitioning including: RANGE, LIST, HASH, COMPOSITE, Automatic LIST	Table Partitioning including: RANGE, LIST	Yes*
Link	Exchange & Split Partitions	N/A	None
Link	Temporary Tables	Temporary Tables	Yes*
Link	Unused Columns	ALTER TABLE DROP COLUMN	Yes
Link	Virtual Columns	Views and/or Function as a Column	Yes*
Link	User Defined Types (UDTs)	User Defined Types (UDTs)	Yes
Link	Read Only Tables & Table Partitions	Read Only Roles and/or Triggers	Yes*
Link	Index Type	Recovery Manager (RMAN)	AWS Aurora Snapshots
Link	B-Tree Index	Flashback Database	AWS Aurora Snapshots
Link	Composite Index	12c Multi-tenant architecture: PDBs and CDB	Databases
Link	BITMAP Index	Tablespaces & DataFiles	Tablespaces
Link	Function-based Index	Data Pump	pg_dump & pg_restore
Link	Global Index	Resource Manager	Separate AWS Aurora Clusters
Link	Identity Column	Database Users	Database Roles
Link	MVCC (Table & Function)	Database Roles	Database Roles
Link	Character Set	SGA & PGA Memory	Memory Buffers
Link	Transaction	V\$ Views & Data Dictionary	System Catalog Tables, Statistics Collector, AWS Aurora Performance Insights
Link		Log Miner	Logging Options
Link		Instance & Database Parameters (SPFILE)	AWS Aurora Parameter Groups
Link		Session Parameters	Session Parameters
Link		Alert.log (error log)	Error Log via AWS Console
Link		Automatic and Manual Statistics Collection	Automatic and Manual Statistics Collection
Link		Viewing Execution Plans	Viewing Execution Plans

<https://aws.amazon.com/dms/getting-started/>

Customer / Partner References

>50,000 Databases Migrated with DMS



Expedia migrated from SQL Server to AWS

Needed real-time
market pricing

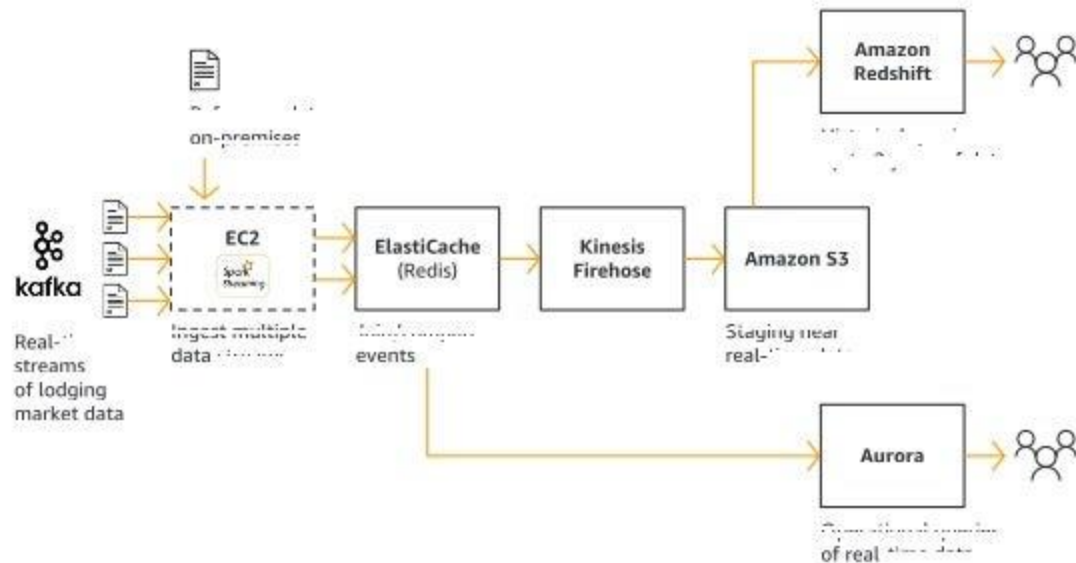
Migrated from Microsoft SQL Server

Use Amazon Aurora, Amazon Redshift,
Kinesis, and ElastiCache

Process high-volume pricing and
availability data

Query execution times reduced
80%–95%

Database has >15B rows and
continues to grow



Trimble migrated from Oracle to Amazon RDS PostgreSQL

"The SCT Assessment Report was the key enabler to allow us to understand the scope of effort required to complete an Oracle to PostgreSQL migration."

What was originally thought to be a largely manual task that no one was particularly excited about having to do became a very straight forward quick and easy process."

- Todd Hofert, Director of Infrastructure Operations, Trimble



Next Steps

Next Steps

- Talk to your AWS account team and AWS Partner
- Ask us about funding for POCs and commercial DB migrations (e.g. Oracle Database to Aurora)
- Read Documentation, White Papers, Playbooks
- Links:
 - DMS & SCT: <https://aws.amazon.com/dms/>
 - Getting Started Guides and Playbooks: <https://aws.amazon.com/dms/getting-started/>

Thank you!



ข้อมูลเชิงลึกของลูกค้าและ Machine Learning (Level 200 -300): การสร้าง Data Lake แบบไร้เซิร์ฟเวอร์บน AWS

Surawut Phornthabthong

Solutions Architect, Amazon Web Services



Architecting a Serverless Data Lake on AWS

Surawut Phornthabthong
Solutions Architect, Amazon Web Services

What is a Data Lake ?

A data lake is an **architectural approach** that allows you to store **massive amounts** of data into a **central location**, so it's readily available to be categorized, processed, analyzed and **consumed by diverse group of users** within an organization.



Challenges faced by data teams

Exponential growth in data



Transactions



Billing



ERP



Web Logs



Sensor Data



Infrastructure Logs



Social

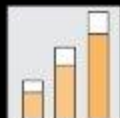
Diversified consumers



Data Scientists



Applications



Business Analyst

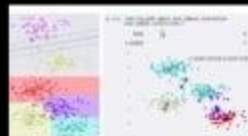


External Consumers

Multiple access mechanisms



API Access



Notebooks



BI Tools

Characteristics of a data lake



Collect
Anything



Dive in
Anywhere



Flexible
Access



Future
Proof

Let's take an example



Device Data

Record-level data

Design Outcomes

Serverless

Ingest & **store** data in real-time

Discover and **catalog** data stored in the lake

Enable batch & real-time **processing**

Consume raw & processed data

Scalable, Highly Available, Pay what you use

Data Lake Ingestion

Multiple data lake ingestion methods



AWS Snowball and AWS Snowmobile

- PB-scale migration



Amazon Kinesis Firehose

- Ingest device streams
- Transform and store on Amazon S3



AWS Storage Gateway

- Migrate legacy files



AWS Direct Connect

- On-premises integration



Native/ISV Connectors

- Ecosystem integration



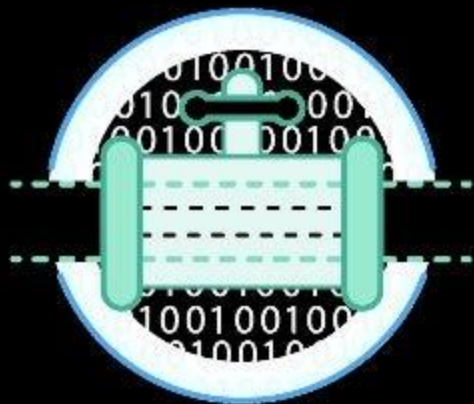
Amazon S3 Transfer Acceleration

- Long-distance data transfer

Amazon Kinesis - real-time analytics



Easily collect, process, and analyze video and data streams in real time



Kinesis Data Streams

Build custom applications that analyze data streams



Kinesis Data Firehose

Load data streams into AWS data stores



Kinesis Data Analytics

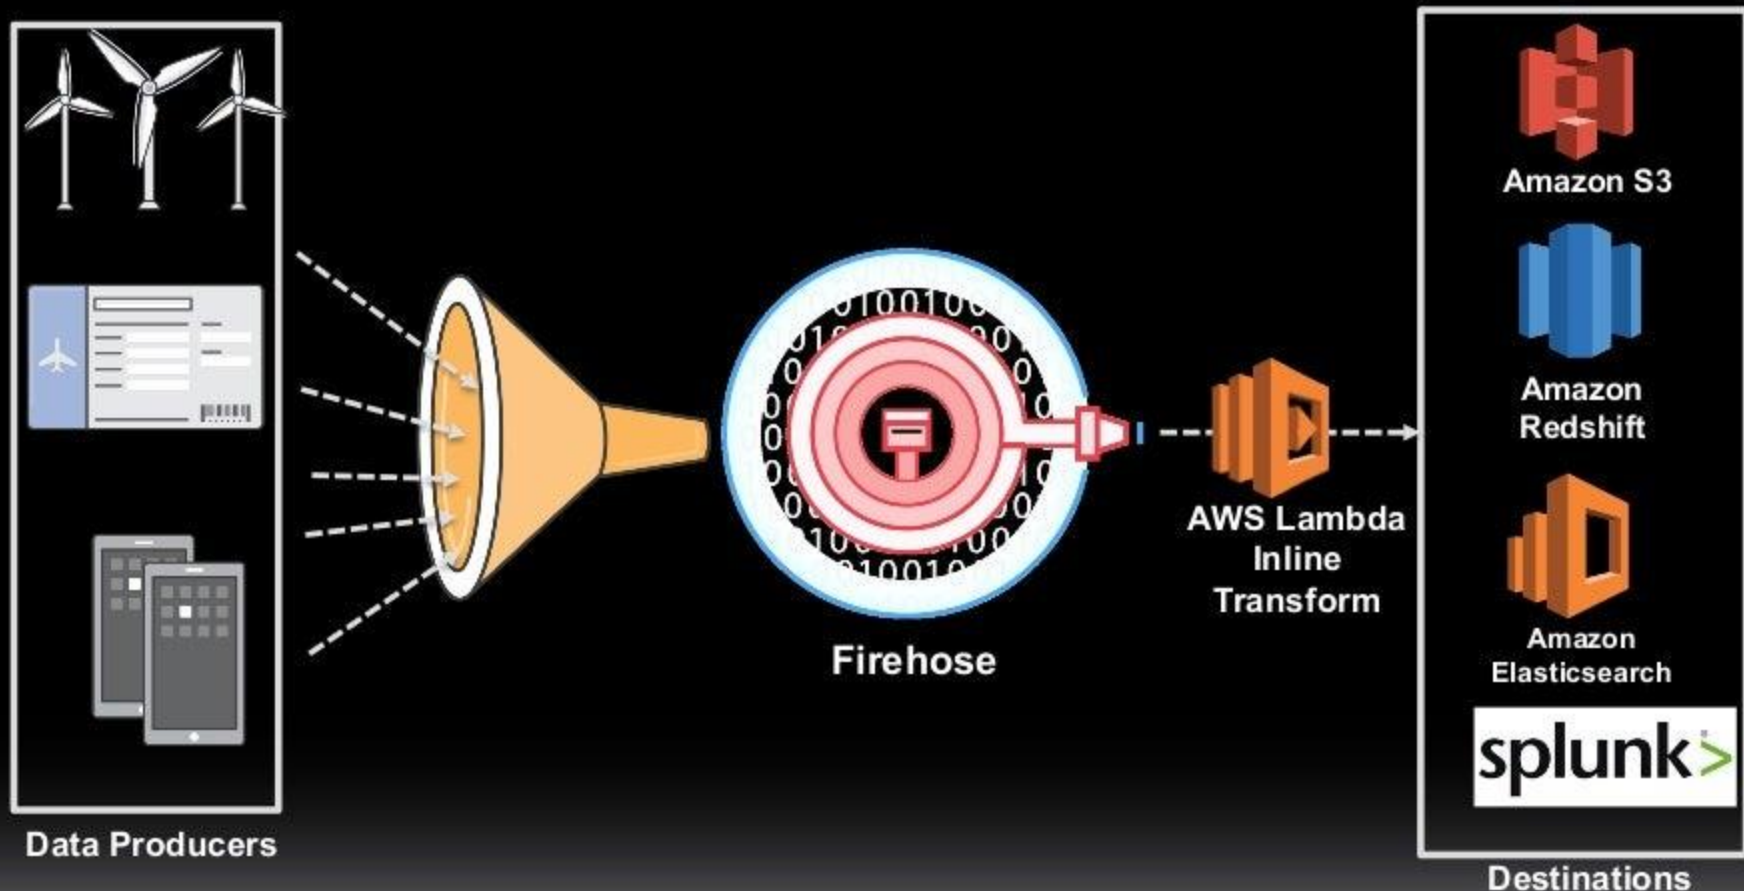
Analyze data streams with SQL



Kinesis Video Streams

Capture, process, and store video streams for analytics

Serverless data delivery with Kinesis Firehose



Architecture



Data Lake Storage

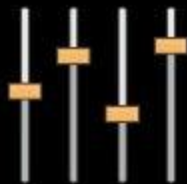
Amazon S3 - Infinite, Durable & Cost Effective Storage



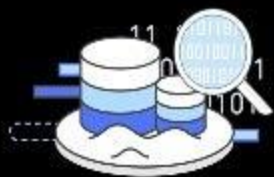
Unmatched durability,
availability, and scalability



Best security, compliance, and audit
capability



Object-level control
at any scale



Business insight into
your data



Most ways to bring
data in



Twice as many partner
integrations

Architecture

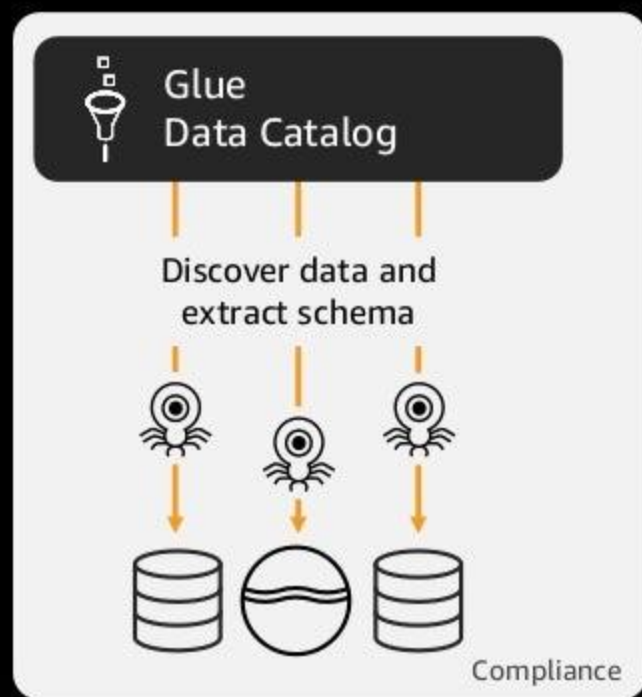


Data Lake Metadata Management

Discover, Catalog & ETL

AWS Glue—data catalog

Make data discoverable



Automatically **discovers data** and **stores schema**

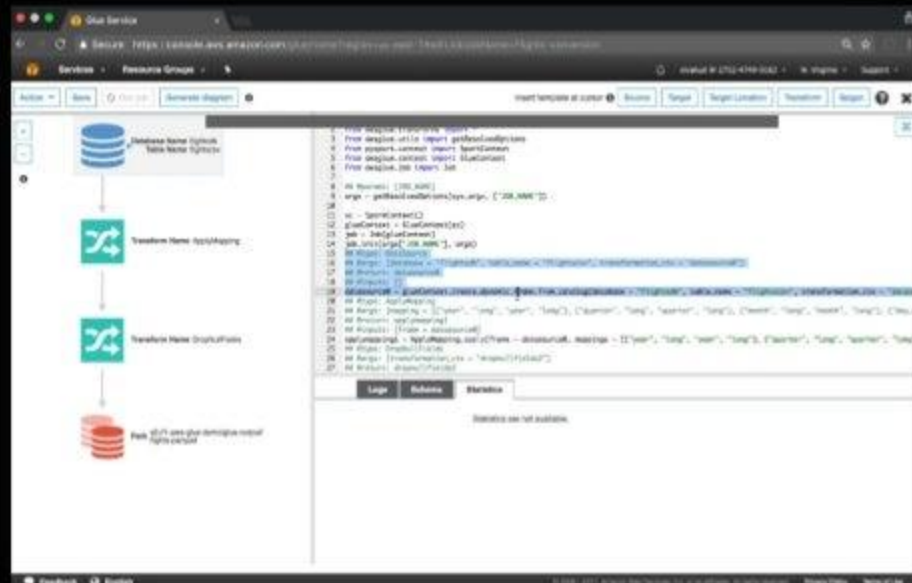
Catalog makes **data searchable**, and available for ETL

Catalog contains table and job definitions

Computes statistics to make queries efficient

AWS Glue—ETL service

Make ETL scripting and deployment easy



Serverless Transformations

Based on **Apache Spark**

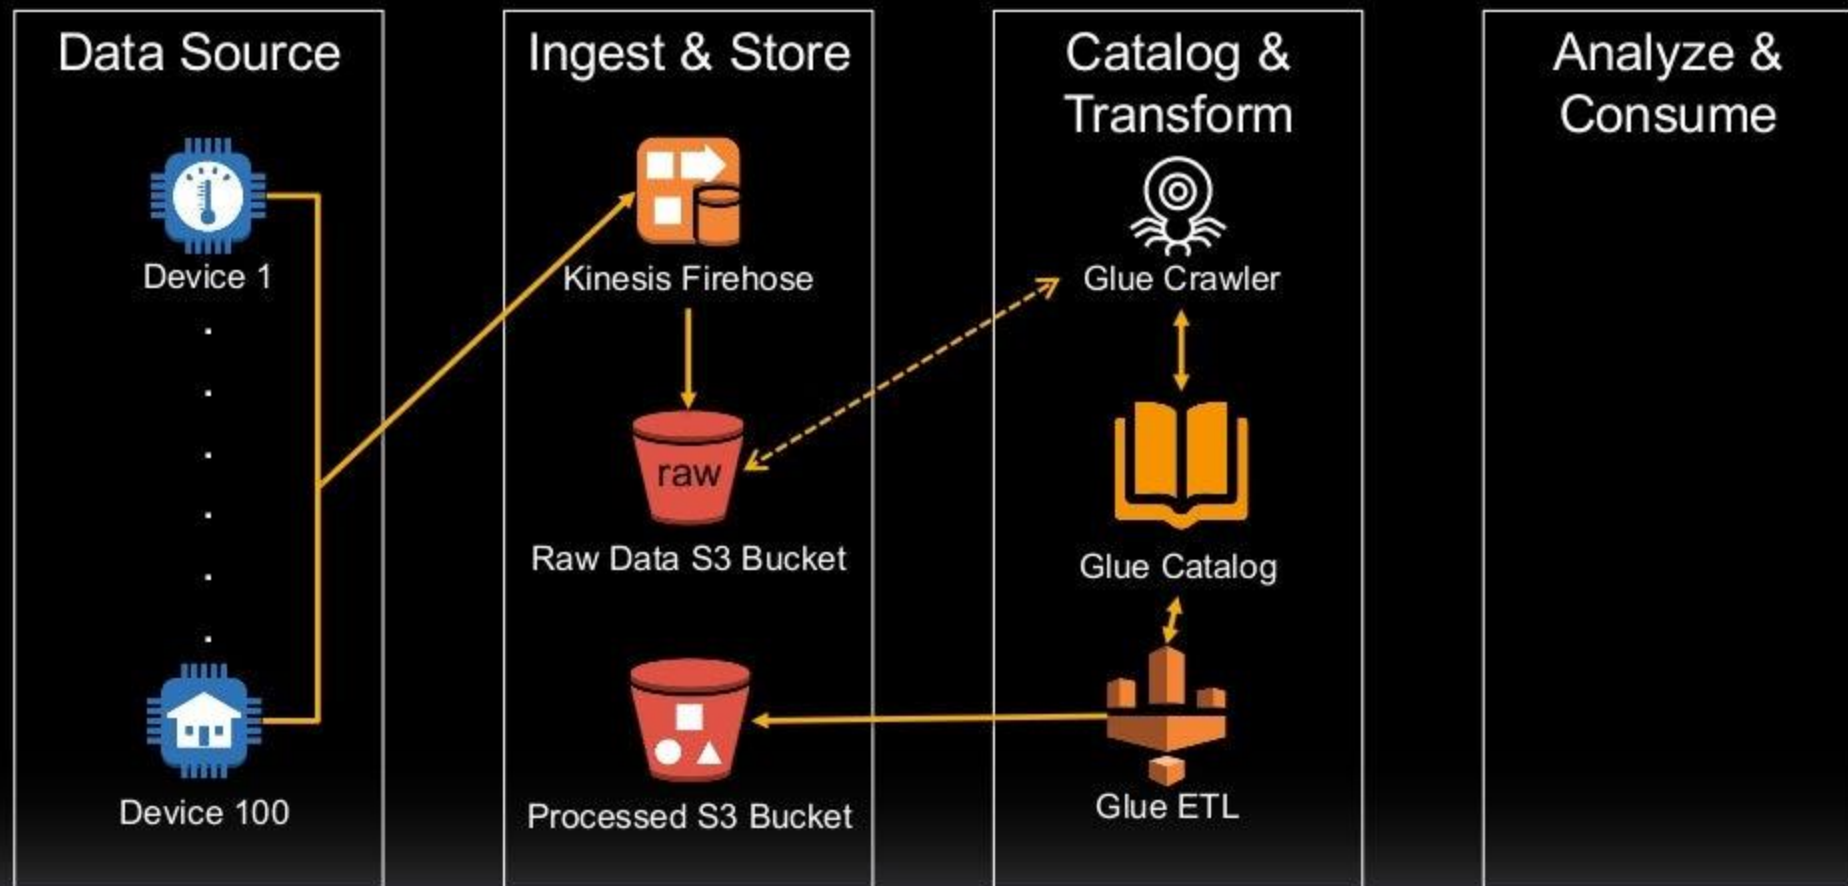
Automatically **generates ETL code**

Code is customizable with **PySpark** and **Scala**

Endpoints provided to edit, debug, test code

Jobs are **scheduled** or **event-based**

Architecture



Data Lake Analytics

Amazon EMR - Big Data Processing

aws SUMMIT

Fully managed – Hadoop Framework

19 Apps : Hadoop, Hive , **Spark**, HBase, **Presto**, and more

S3 Integration – **Decouple Compute and Storage**

Low Cost – **Transient Clusters**, Per Second Pricing, Spot Instances



Amazon Redshift - Modern Data Warehousing

Fast, scalable, **fully managed EDW** at **1/10th the cost** of other EDWs

Massively parallel, scales from gigabytes to **exabytes**

Access data across your Redshift DW and Amazon S3 **data lake**



Amazon Athena—Interactive Analysis



Interactive query service to analyze data in Amazon S3 using standard SQL

No infrastructure to set up or manage and no data to load

Query Instantly



Zero setup cost; just point to S3 and start querying

Pay per query



Pay only for queries run; save 30–90% on per-query costs through compression

Open



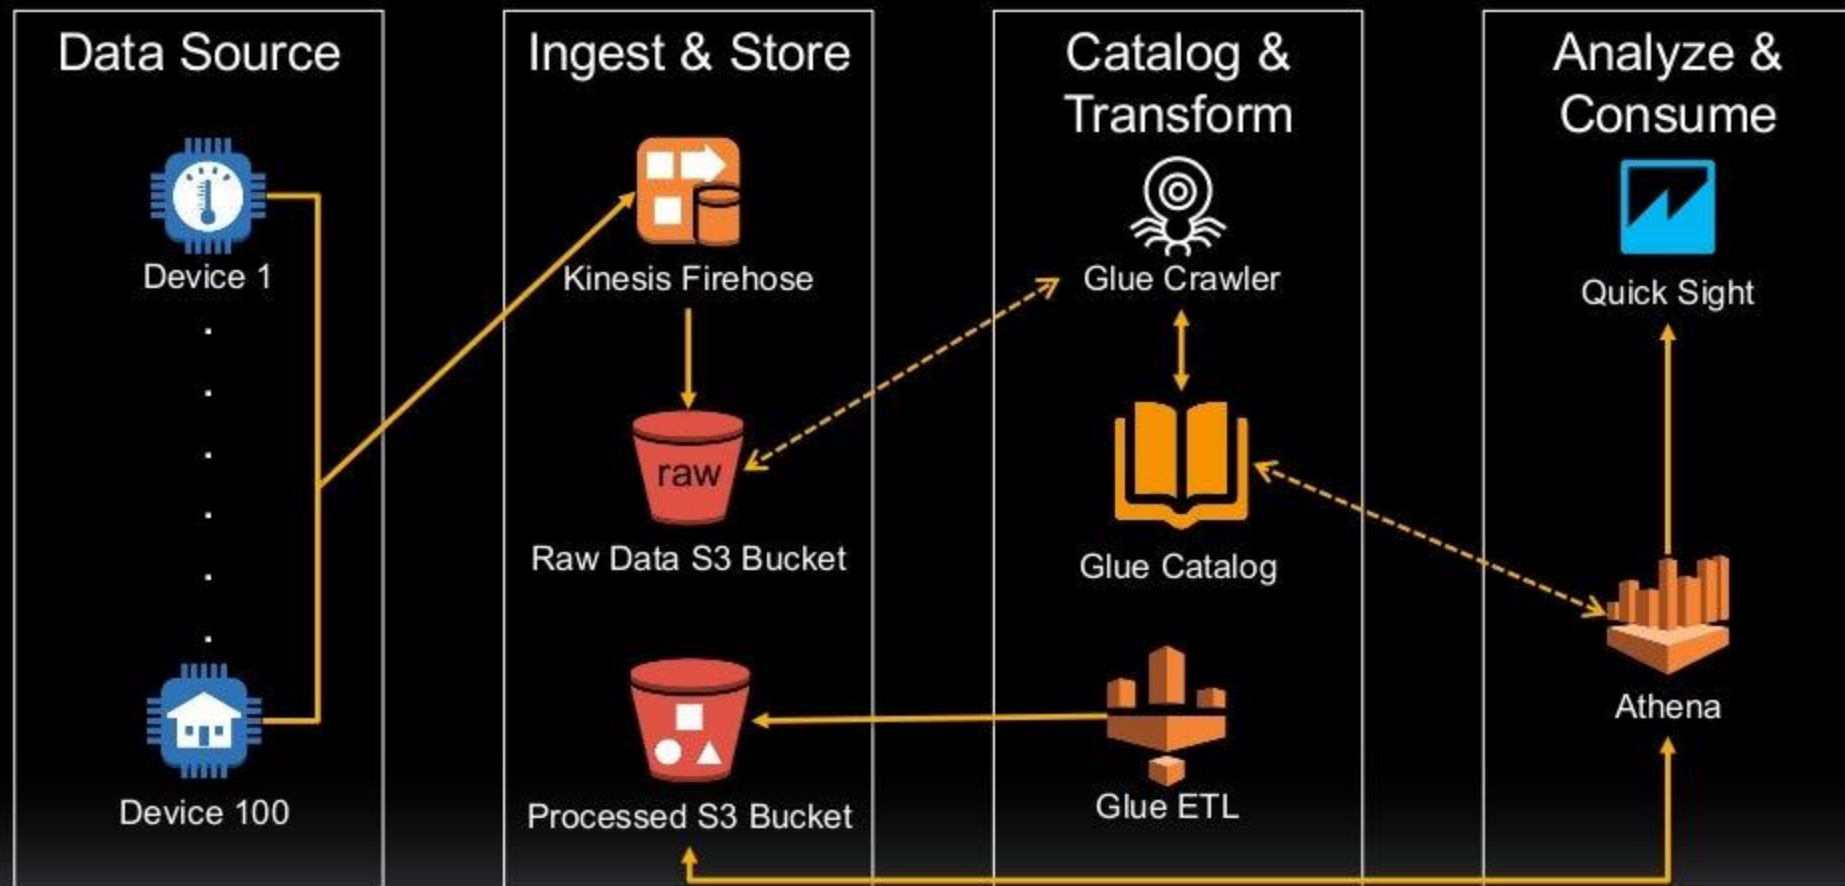
ANSI SQL interface, JDBC/ODBC drivers, multiple formats, compression types, and complex joins and data types

Easy



Serverless: zero infrastructure, zero administration
Integrated with QuickSight

Architecture



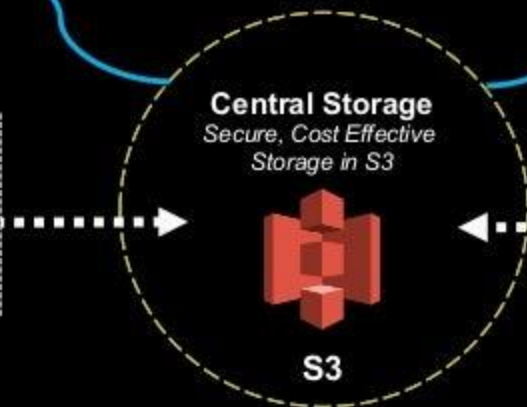
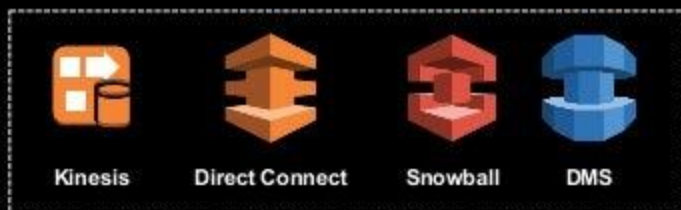
Take the demo home...



<http://bit.ly/sg-summit-datalake-demo>



Data Ingestion
Get your data into S3 quickly and securely



Processing & Analytics
Use predictive and prescriptive analytics to gain better understanding





Daniel Muller

Head of Cloud Infrastructure, Spuul



danielmullerch





Watch unlimited
Bollywood Movies.
Anytime. Anywhere.

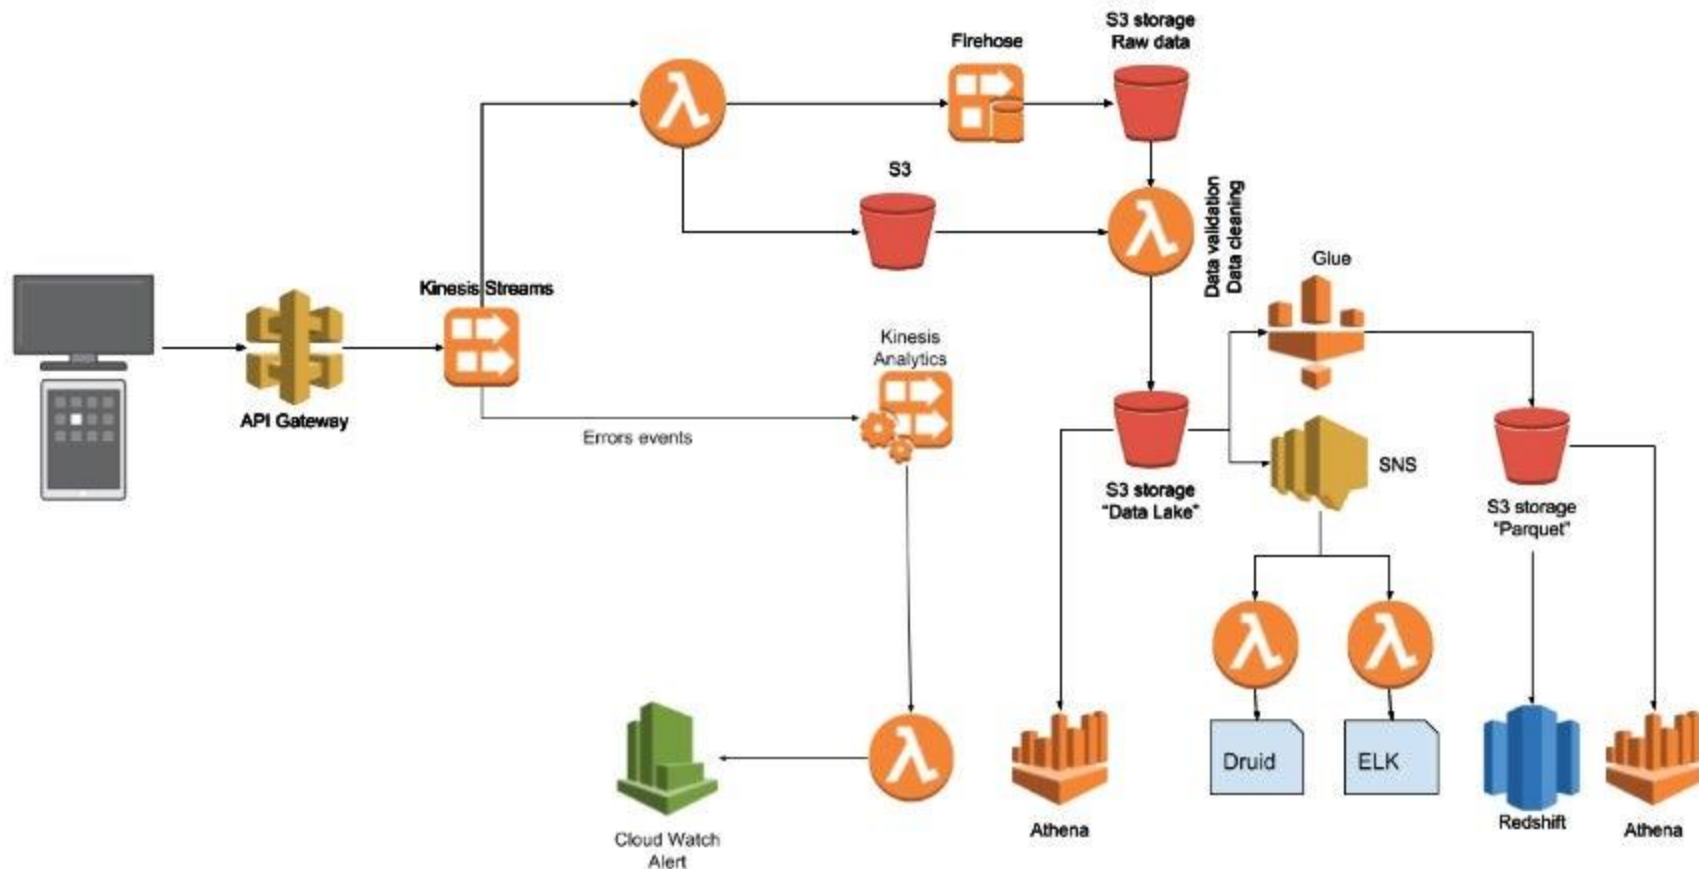


- ❖ Leading OTT player
- ❖ Indian Movies, Shows and Live TV
- ❖ 50 millions registered users
- ❖ Users in 5 continents
- ❖ Content served on Mobile, connected TVs, SetTop Boxes
- ❖ (Coming Soon..!) Non-Indian content

Why we built a serverless data lake ?

- **100+ event types** - across microservices & devices
- **Flexibility** - Ingestion, Consumption
- **Bottomless storage** - Cheap & Reliable
- **Ad-hoc querying** - Analyze data without a data-warehouse
- **Future use cases** - 3rd party integrations
- We **hate managing servers**..! #NoMoreServers

Architecture



Lessons Learnt & Best Practices

Use a framework - SAM, Serverless, Apex/Up

Store data in raw format - debugging and re-processing

Convert to Columnar Formats - Optimized for reads

Partition data - Based on your filters

Specify columns to load - Reduce data transfer

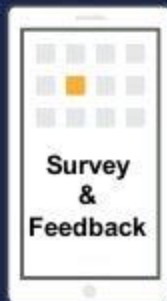
Create files of ~100MB - Reduces S3 list API calls

Compress Data in Lake - Reduce network transfers

Use Lambda for Automation - Wire things together

Summary

Thank You !



Take the demo home
<http://bit.ly/sg-summit-datalake-demo>

#AWSSummitSG

#NoMoreServers

#DataLake

 **unni_k_pillai**

Thank You

You will receive today's webinar recording and presentation deck,
look out for it in your inbox.